

CS 537: Introduction to Operating Systems

Spring 2024: Final Exam

This exam is closed book, closed notes.

No calculators may be used. All cell phones must be turned off and put away.

You have 2 hours to complete this exam.

Write all of your answers on the accu-scan form with a #2 pencil:

- CANVAS LAST NAME - fill in your last (family) name starting at the leftmost column
- CANVAS FIRST NAME - fill in the first five letters of your first (given) name
- IDENTIFICATION NUMBER - This is your UW Student WisCard ID number
- ABC of SPECIAL CODES - Write your lecture number as a three digit value:
 - 001 (Afternoon Lecture)
 - 002 (Morning Lecture)
- **D of SPECIAL CODES - Write the value 1 for this version of the exam.**

These exam questions must be returned at the end of the exam, but we will not grade anything in this booklet.

You may separate the pages of this exam if it helps you.

Unless stated (or implied) otherwise, you should make the following assumptions:

- The OS manages a single uniprocessor (single core)
- All memory is byte addressable
- The terminology lg means \log_2
- 2^{10} bytes = 1KB
- 2^{20} bytes = 1MB
- Data is allocated with optimal alignment, starting at the beginning of a page
- Assume leading zeros can be removed from numbers (e.g., 0x06 == 0x6).
- Hex numbers are represented with a proceeding "0x"

This exam has 75 questions. Each question has the same number of points.

Good luck!

Part 1. Virtualization

True / False

1. If a PTE is invalid, accessing the corresponding virtual address will trigger a segmentation fault.
2. A preemptive scheduler relies on timer interrupts to force processes to yield the CPU.
3. Processes running the same program can share memory for the code using virtual memory.
4. Since STCF is the ideal scheduler, it is the basis for modern operating system schedulers.
5. The cost of a context switch is proportional to the number of page table entries in the new process, since they have to be restored.

Multiple Choice

Address Translation -- You have a system that uses a linear page table for virtual to physical address translation. The system has the following parameters:

Address Space size: 16 KB
Physical Memory size: 64KB
Page size: 4KB

The format of the page table has the high-order (left-most) bit as the VALID bit. If the bit is 1 then the rest of the entry is the PFN. If the bit is 0 then the page is not valid. The current page table (from entry 0 down to the last entry) is:

0x80000000
0x00000000
0x00000000
0x8000000a

For each virtual address, indicate the physical address it translates to or that it is an invalid address.

6. Virtual Address 0x00001b00 (decimal: 6912)
A. 0x00000b00 (decimal: 2816)
B. 0x00000000 (decimal: 0)
C. 0x0000b00a (decimal: 45066)
D. 0x0000ab00 (decimal: 43776)
E. Invalid Address
7. Virtual Address 0x00003919 (decimal: 14617)
A. 0x00000919 (decimal: 2329)
B. 0x00000519 (decimal: 1305)
C. 0x00003919 (decimal: 14617)
D. 0x0000a919 (decimal: 43289)
E. Invalid Address

CPU Job Scheduling -- A computer has the following workload. If needed, assume a time-slice of 1 second and **if there is a choice the newest arriving job is selected.**

Job Name	Arrival Time (seconds)	CPU Burst Time (seconds)
A	0	8
B	4	4
C	7	6

8. Given a FIFO scheduler, what is the turnaround time of job B?

- A. 4
- B. 6
- C. 7
- D. 8
- E. None of the above

9. Given a RR scheduler, what is the average response time?

- A. 0
- B. 2
- C. 6
- D. 7
- E. None of the above

10. Given a STCF scheduler, what is the average turnaround time?

- A. 7
- B. 9
- C. 14
- D. 18
- E. None of the above

Part 2. Concurrency

True / False

11. A spinlock is a good choice of lock to protect a critical section that takes 2-3 cycles.

12. If two threads simultaneously (that is, without synchronization) increment a memory address by 1 with a load/add/store assembly instruction sequence, the result can be an increment by either 1 or 2 (with no other possibilities).

13. The primary purpose of mutex locks is to make sure that one thread runs before another.

14. Semaphores can be used to replace condition variables.

15. A semaphore used to create mutual exclusion for critical sections should be initialized to 1.

Multiple Choice

16. Which one of the following is NOT a necessary condition for deadlock to occur?

- A. mutual exclusion
- B. hold-and-wait
- C. no preemption
- D. circular wait
- E. All the above are necessary

Consider the following producer-consumer code. The producer thread runs for 'loops' number of iterations and is filling a buffer of size 'MAX'. The buffer holds 'count' items. Consider one producer and one consumer thread.

```
void *producer(void *arg) {
    int i;
    for (i=0; i<loops; i++) {
        mutex_lock(&mutex);          //p1
        while (count == MAX)         //p2
            cond_wait(&empty,&mutex); //p3
        put(i);                      //p4
        cond_signal(&fill);           //p5
        mutex_unlock(&mutex);         //p6
    }
}
```

```
void *consumer(void *arg) {
    int i;
    for (i=0; i<loops; i++) {
        mutex_lock(&mutex);          //c1
        while (count == 0)           //c2
            cond_wait(&fill,&mutex); //c3
        int tmp = get();              //c4
        cond_signal(&empty);          //c5
        mutex_unlock(&mutex);         //c6
        printf("%d\n",tmp);           //c7
    }
}
```

17. What are the sequence of lines executed when the producer runs for one iteration followed by the consumer? Assume you start from i=0, count=0, and MAX=5.

- A. p1 p2 p3 p4 p5 p6 c1 c2 c3 c4 c5 c6 c7
- B. p1 p2 p3 p4 p5 p6 c1 c2 c4 c5 c6 c7
- C. p1 p2 p4 p5 p6 c1 c2 c4 c5 c6 c7
- D. p1 p2 p4 p5 p6 c1 c2 c3 c4 c5 c6 c7
- E. None of the above

18. What are the sequence of lines executed when the consumer thread runs first?
- A. c1 c2 c3 p1 p2 p4 p5 p6 c2 c4 c5 c6 c7
 - B. c1 p1 p2 p4 p5 p6 c2 c4 c5 c6 c7
 - C. c1 c2 c3 p1 p2 p3 p4 p5 p6 c4 c5 c6 c7
 - D. c1 c2 c4 c5 c6 c7 p1 p2 p4 p5 p6
 - E. None of the above
19. This is a correct implementation of the producer/consumer problem.
- A. True
 - B. False
20. What is the objective of a reader/writer lock?
- A. Only one thread of any type can hold the lock at a time
 - B. Either only one reader thread can hold the lock at a time or multiple writer threads can hold the lock at the same time
 - C. Either only one writer thread can hold the lock at a time or multiple reader threads can hold the lock at the same time
 - D. Multiple reader or writer threads can hold the lock at the same time
 - E. None of the above

Part 3. Persistence

True/False

21. Because of the physical nature of SSDs they only efficiently support log-structured file systems.
22. In LFS the location of an inode is not fixed and so another level of indirection is used to map an inode number to its location on disk.
23. The reason why data journaling waits to write the transaction commit to the journal until all other writes for the transaction to the journal have completed is so that a crash can be detected when writing to the data portion of the file system.
24. When fsck discovers an inconsistency between the inode bitmap and the inode table, it is resolved by trusting the information within the inodes.
25. The fast file system (FFS) improved performance by keeping inodes close to the data that they referred to and by keeping files in the same directory close to one another.
26. When a process reads data from the end of a file, the bitmap is consulted to check if the blocks have been allocated.
27. The system call to open `"/foo/bar/blah.txt"` requires exactly 5 read accesses to disk. Assume each inode access is one read, and all directories fit into one disk block.
28. The directory data contains mappings between file names and the associated inode number and this information is stored in the data block region of the file system.

29. Extents are an alternative to individual block pointers to data in an inode.
30. Assuming a block size of 4KB and 4-byte block pointers, an inode with 12 direct pointers and 1 indirect pointer supports a maximum file size of just over 6MB.
31. When an indirect pointer is used in an inode, it points to a block of other inodes in the inode table.
32. The advantage of using a free list rather than a data bitmap for tracking free blocks in a file system is that no search is needed to allocate a block.
33. The execute permission bit for a directory enables a user (or group or everyone) to do things like navigate into the given directory.
34. One reason that symbolic links cannot refer to a directory is because of the fear of creating cycles in the directory tree.
35. Only when the reference count of a file reaches zero does the file system free the inode and related data blocks.
36. The metadata for a regular file inode includes the number of hard links to the file.
37. Opening the same file twice in a process without closing it first will result in multiple references to the same file struct in the open file table.
38. The file descriptor, otherwise known as the inode number, is the handle used in Unix systems to access files.
39. The one main difference between RAID levels 4 and 5 is where the parity information is stored.
40. Of the RAID levels covered in class, the two that maximize capacity while still providing some type of redundancy are RAID levels 0 (striped) and 5 (rotating parity).
41. The SCAN algorithm is optimal in maximizing the throughput rate of a hard disk.
42. The purpose of track-skew is to make sure that sequential reads/writes crossing track boundaries are serviced optimally, not requiring an additional rotation.
43. In a write-back caching disk a write is acknowledged as soon as the data is placed in the cache as opposed to write-through caching.
44. If we create a symbolic link to a file, the reference count for that file's inode is not incremented
45. A log-based file system attempts to make all writes sequential.
46. For a 15,000 RPM disk, the average rotation time for a seek is 4ms
47. A device driver is software that executes on the microcontroller of a peripheral device.

48. Memory-mapped IO and explicit IO instructions are two alternate architectural approaches for how an OS accesses device registers.

49. Programmed I/O was introduced to relieve the OS of transferring large chunks of data to and from a device's registers for it to perform an IO operation.

50. Polling wastes less CPU time compared to interrupts to determine when a disk drive has finished an IO operation.

51. Over 70% of the Linux kernel code is found in device drivers.

Multiple Choice

An Operating System can perform the following IO Operations:

mkdir() - create a new directory

creat() - create a new, empty file

open(), write(), close() - appends a block to a file (this is one operation)

link(path1, path2) - creates a hard link at path2 to a file (path1)

unlink() - unlinks a file (removing it if refcnt==0)

A file system like the VSFS discussed in class is used on disk, which has an inode bitmap, a data bitmap, an inode table, and data blocks. Both inodes and data blocks are indexed starting at 0. An inode has three fields: type (f for file and d for directory), address of data block (either -1 if no data or a single integer address for the data block), and a reference count saying how many directory entries there are to this inode. Data blocks are indicated by matching square brackets and can contain file data (a single character) or directory information (name-inode pairs). Directory data always fits into a single data block.

Here is the initial state of the file system:

```
inode bitmap  10000000
inodes        [d a:0 r:2] [] [] [] [] [] [] []
data bitmap   10000000
data          [(.,0) (.,0)] [] [] [] [] [] []
```

52. What IO operation took place to change the state of the file system to the following?

```
inode bitmap  11000000
inodes        [d a:0 r:3][d a:1 r:2][][][][][][]
data bitmap   11000000
data          [(.,0) (.,0) (k,1)][(.,1) (.,0)][][][][][]
```

A. mkdir("/k")

B. creat("/k")

C. link("/", "/k")

D. unlink("/k")

E. None of the above or Can not determine

53. What IO operation took place to continue changing the state of the file system to the following (from the previous question)?

```
inode bitmap 11100000
inodes       [d a:0 r:3][d a:1 r:2][f a:-1 r:1][ ][ ][ ][ ][ ]
data bitmap  11000000
data         [(.,0) (.,0) (k,1)][(.,1) (.,0) (h,2)][ ][ ][ ][ ][ ]
```

- A. creat("/h")
- B. creat("/k/h")
- C. link("/k","/h")
- D. link("/k","/k/h")
- E. None of the above or Can not determine

54. What IO operation took place to continue changing the state of the file system to the following (from the previous question)?

```
inode bitmap 11100000
inodes       [d a:0 r:3][d a:1 r:2][f a:-1 r:2][ ][ ][ ][ ][ ]
data bitmap  11000000
data         [(.,0) (.,0) (k,1) (h,2)][(.,1) (.,0) (h,2)][ ][ ][ ][ ][ ]
```

- A. link("/k/h","/h")
- B. link("/k","/h")
- C. open("/k/h"),write("h"),close()
- D. open("/k"),write("h"),close()
- E. None of the above or Can not determine

55. What IO Operation took place to continue changing the state of the file system to the following (from the previous question)?

```
inode bitmap 11100000
inodes       [d a:0 r:3][d a:1 r:2][f a:2 r:2][ ][ ][ ][ ][ ]
data bitmap  11100000
data         [(.,0) (.,0) (k,1) (h,2)][(.,1) (.,0) (h,2)][u][ ][ ][ ][ ]
```

- A. open("/k/h"),write("u"),close()
- B. open("/h"),write("u"),close()
- C. open("/k"),write("u"),close()
- D. creat("/u")
- E. None of the above or Can not determine

56. What 2 IO operations took place to continue changing the state of the file system to the following (from the previous question)?

```
inode bitmap 11100000
inodes       [d a:0 r:3][d a:1 r:2][f a:2 r:2][ ][ ][ ][ ][ ]
data bitmap  11100000
data         [(.,0) (.,0) (k,1) (h,2)][(.,1) (.,0) (s,2)][u][ ][ ][ ][ ]
```

A. creat("/k/s")
unlink("/k/h")

B. link("/k/h","/k/s")
unlink("/k/h")

C. unlink("/k/h")
link("/h","/k/s")

D. unlink("/k/h")
mkdir("/k/s")

E. None of the above or Can not determine

This same file system may contain errors or inconsistencies in its data structures. For each of the following file system states, identify the error or inconsistency.

57.

```
inode bitmap 11100000
inodes       [d a:0 r:3][d a:1 r:2][f a:2 r:2][ ][ ][ ][ ][ ]
data bitmap  11100000
data         [(.,0) (.,0) (k,1) (h,2)][(.,0) (.,1) (h,2)][u][ ][ ][ ][ ]
```

- A. Inode bitmap mismatch with inode table
- B. Inode's refcount is incorrect
- C. Directory's data connecting it to the directory tree is incorrect
- D. File system is in a consistent state
- E. None of the above

58.

```
inode bitmap 11100000
inodes       [d a:0 r:3][d a:1 r:3][f a:2 r:2][ ][ ][ ][ ][ ]
data bitmap  11100000
data         [(.,0) (.,0) (k,1) (h,2)][(.,1) (.,0) (h,2)][u][ ][ ][ ][ ]
```

- A. Inode bitmap mismatch with inode table
- B. Inode's refcount is incorrect
- C. Directory's data connecting it to the directory tree is incorrect
- D. File system is in a consistent state
- E. None of the above

59.

```
inode bitmap 11000000
inodes       [d a:0 r:3][d a:1 r:2][f a:2 r:2][ ][ ][ ][ ][ ][ ]
data bitmap  11100000
data         [(.,0) (.,0) (k,1) (h,2)][(.,1) (.,0) (h,2)][u][ ][ ][ ][ ][ ]
```

- A. Inode bitmap mismatch with inode table
- B. Inode's refcount is incorrect
- C. Directory's data connecting it to the directory tree is incorrect
- D. File system is in a consistent state
- E. None of the above

A journaling system is added to the above file system to prevent inconsistencies. For each of the following IO operations and journal types, of the options listed, identify the one that is written to the journal.

60. creat("/foo") using data journaling.

- A. inode bitmap
- B. data bitmap
- C. a new data block
- D. More than one of the above are written to the journal
- E. None of the above are written to the journal

61. mkdir("/bar") using metadata journaling.

- A. data bitmap
- B. root's directory data
- C. a new data block
- D. More than one of the above are written to the journal
- E. None of the above are written to the journal

62. open("/foo"),write("x"),close() using metadata journaling to write to an empty file.

- A. root's directory data
- B. a new data block
- C. inode bitmap
- D. More than one of the above are written to the journal
- E. None of the above are written to the journal

63. link("/foo","foo2") using data journaling.

- A. inode bitmap
- B. root's directory data
- C. a new data block
- D. More than one of the above are written to the journal
- E. None of the above are written to the journal

You decide to change the file system to a log-structured file system. In this file system, each block may contain one of the following:

- * A checkpoint region containing block addresses of the live chunks of the imap
- * An inode containing its type (regular file or dir), size (in blocks), reference count, and pointers to data blocks
- * Directory data (name to inode number mappings)
- * A chunk of the imap listing blocks containing inodes (the position (index) within the chunk is the inode number)
- * File data

INITIAL file system contents:

```
[ 0 ] live checkpoint: 3 -- -- -- -- -- -- -- -- -- -- -- -- -- -- -- --
[ 1 ] live [.,0] [...,0] -- -- -- -- -- -- -- --
[ 2 ] live type:dir size:1 refs:2 ptrs: 1 -- -- -- -- -- -- -- --
[ 3 ] live chunk(imap): 2 -- -- -- -- -- -- -- -- -- -- -- -- -- -- -- --
```

Notice that the initial file system consists of 4 live blocks.

As the file system is updated by IO operations, you will be presented with the contents of the file system and will be asked to identify which blocks of the file system are live and what IO operations were performed.

64. Three IO operations were performed on the initial file system state to reach the state shown below. Which blocks are live?

```
[ 0 ] checkpoint: 16 -- -- -- -- -- -- -- -- -- -- -- -- -- -- -- --
[ 1 ] [.,0] [...,0] -- -- -- -- -- -- -- --
[ 2 ] type:dir size:1 refs:2 ptrs: 1 -- -- -- -- -- -- -- --
[ 3 ] chunk(imap): 2 -- -- -- -- -- -- -- -- -- -- -- -- -- -- -- --
[ 4 ] [.,0] [...,0] [my4,1] -- -- -- -- -- -- -- --
[ 5 ] type:dir size:1 refs:2 ptrs: 4 -- -- -- -- -- -- -- --
[ 6 ] type:reg size:0 refs:1 ptrs: -- -- -- -- -- -- -- -- -- --
[ 7 ] chunk(imap): 5 6 -- -- -- -- -- -- -- -- -- -- -- -- -- -- -- --
[ 8 ] [.,0] [...,0] [my4,1] [ke3,1] -- -- -- -- -- -- -- --
[ 9 ] type:dir size:1 refs:2 ptrs: 8 -- -- -- -- -- -- -- --
[10] type:reg size:0 refs:2 ptrs: -- -- -- -- -- -- -- -- -- --
[11] chunk(imap): 9 10 -- -- -- -- -- -- -- -- -- -- -- -- -- -- -- --
[12] [.,0] [...,0] [my4,1] [ke3,1] [ka0,2] -- -- -- -- -- -- -- --
[13] [.,2] [...,0] -- -- -- -- -- -- -- --
[14] type:dir size:1 refs:3 ptrs: 12 -- -- -- -- -- -- -- --
[15] type:dir size:1 refs:2 ptrs: 13 -- -- -- -- -- -- -- --
[16] chunk(imap): 14 10 15 -- -- -- -- -- -- -- -- -- -- -- -- -- -- -- --
```

- A. 0, 2, 5, 6, 16
- B. 0, 8, 9, 10, 16
- C. 13, 15
- D. 0, 10, 12, 13, 14, 15, 16
- E. None of the above

65. What three IO operations would result in the above file system state?

A. create file /my4
create file /ke3
create file /ka0

B. create file /my4
link file /my4 to /ke3
create dir /ka0

C. create dir /my4
create file /ke3
write data to file /ke3

D. create file /my4
write data to file /my4
create dir /ka0

E. None of the above

66. What block contains the live root inode?

A. 2
B. 5
C. 9
D. 14
E. None of the above

You use an SSD to hold your data. The SSD has 3 blocks with 10 pages per block. Each page holds a single character. The SSD has been used for a while and the current state of the SSD is shown below.

The diagram for the SSD contains the following items:

FTL mapping logical pages to physical pages (1st line)

block number (2nd line)

page number (3rd and 4th lines)

state of each page (valid, Erased, or invalid) (5th line)

data stored at each page (6th line)

an indicator(+) if a page is currently live (has an entry in the FTL) (7th line)

Current State of SSD:

FTL 1:6 4:5 19:7 20:3

Block	0	1	2
Page	0000000000	1111111111	2222222222
	0123456789	0123456789	0123456789
State	vvvvvvvvEE	iiiiiiiiiii	iiiiiiiiiii
Data	z9fFAdux		
Live	+ +++		

The OS can perform 3 operations:

read(logical page #) -- if page is live return the character at the page, otherwise error

write(logical page #, character) -- writes a character to the logical page #.

erase(logical page #) -- removes logical page # from the FTL mapping if there, otherwise error

67. What single OS operation will result in the following SSD state?

```
FTL      1:6    4:8   19:7   20:3

Block 0           1           2
Page  0000000000 111111111 222222222
      0123456789 0123456789 0123456789
State  vvvvvvvvvvE iiiiiiiiii iiiiiiiiii
Data   z9fFAduxz
Live   +   +++
```

- A. write(8,z)
- B. read(8)
- C. write(4,z)
- D. erase(4)
- E. None of the above

68. Continuing making OS operations to the SSD, a read(20) is performed. What is the return value?

- A. z
- B. F
- C. v
- D. error
- E. None of the above

69. Another OS operation is performed which changes the state of the SSD again to that shown below. What was the operation?

```
FTL      1:6    2:9    4:8   19:7    20:3

Block 0           1           2
Page  0000000000 111111111 222222222
      0123456789 0123456789 0123456789
State  vvvvvvvvvv iiiiiiiiii iiiiiiiiii
Data   z9fFAduxza
Live   +   ++++
```

- A. write(9,a)
- B. erase(4)
- C. write(2,z)
- D. write(1,9)
- E. None of the above

70. Now the OS performs a write(28,q) operation. What low-level SSD operations must be performed in order to carry out this OS operation?

A. erase block 1

program q to physical page 10

Enter mapping 28:10 in FTL

B. read physical pages 3, 6, 7, 8, 9

erase block 0

program physical pages F to 10, u to 11, x to 12, z to 13, a to 14, q to 15

C. program q to physical page 10

Enter mapping of 10:28 in FTL

D. erase page 10

program q to physical page 10

Enter mapping 28:10 in FTL

E. None of the above

71. Some time later the state of the SSD is as shown below. The garbage collector decides to reclaim block 0. What will the state of the SSD be after the garbage collector has finished?

```
FTL      1:6    2:9    4:13   5:12
         7:11   19:7   20:3    23:14
         28:10

Block 0      1      2
Page 0000000000 111111111 222222222
      0123456789 0123456789 0123456789
State vvvvvvvvvv vvvvvEEEEEE iiiiiiiiii
Data  z9fFAduxza qh7Et
Live   +  ++ + +++++
```

A.

```
FTL      1: 6    2: 9    4: 13   5: 12
         7: 11   19: 7   20: 3    23: 14
         28: 10

Block 0      1      2
Page 0000000000 111111111 222222222
      0123456789 0123456789 0123456789
State vvvvvvvvvv vvvvvEEEEEE iiiiiiiiii
Data  z9fFAduxza qh7EtFuxa
Live   +  ++ + +++++
```

B.

```
FTL      1:21    2:24    4:13    5:12
          7:11    19:22   20:18   23:14
          28:10

Block 0           1           2
Page  0000000000 111111111 222222222
      0123456789 0123456789 0123456789
State  vvvvvvvvvv vvvvvEEEEEE iiiiiiiiii
Data   z9fFAduxza qh7Etz9fFA duxza
Live   ++++++  +  ++ +
```

C.

```
FTL      1:21    2:24    4:13    5:12
          7:11    19:22   20:18   23:14
          28:10

Block 0           1           2
Page  0000000000 111111111 222222222
      0123456789 0123456789 0123456789
State  EEEEEEEEEEE vvvvvvvvvv vvvvvEEEEEE
Data   qh7Etz9fFA duxza
Live   ++++++  +  ++ +
```

D.

```
FTL      1:16    2:18    4:13    5:12
          7:11   19:17   20:15   23:14
          28:10

Block 0           1           2
Page  0000000000 111111111 222222222
      0123456789 0123456789 0123456789
State  vvvvvvvvvv vvvvvvvvvvE iiiiiiiiii
Data   z9fFAduxza qh7EtFuxa
Live   ++++++++
```

E. None of the above

Part 4. Distributed Systems and NFS

72. With the TCP protocol, what does a receiver do if it receives a packet that it has already received?

- A. Drop the packet
- B. Send an acknowledgement and drop the packet
- C. Send an acknowledgement and forward the packet to the application
- D. Forward the packet to the application
- E. None of the above

73. What reliability is built into the UDP protocol?

- A. Packet loss detection
- B. Mechanism for detecting data corruption
- C. Guarantees data is received in order it was sent
- D. Mechanism for sender to execute code on the receiver
- E. None of the above

74. In NFS, a client opens a file and obtains file descriptor 145 and a file handle <volume = 1, inode = 23, generation = 2>. After a few seconds, the original client issues a request for this file and the server finds the generation number is 3 and returns an error to the client. What must have happened in the meantime?

- A. The server crashed and rebooted
- B. File descriptor 145 was closed and reopened with a different file
- C. The file for inode 23 was removed and re-allocated for a different file/directory
- D. Another client wrote to this file and the original client needs to update its local cached version
- E. This situation is impossible, there is a bug in the server

75. In the NFS protocol, the mkdir operation is not quite idempotent. How does this affect clients?

- A. No effect - it only affects performance on an unreliable network
- B. The operation can succeed even if the directory already exists
- C. The operation can fail and do nothing even though the directory does not exist
- D. The operation can fail even though it successfully created the directory
- E. None of the above

That's it. Thank you for taking CS 537 and have a great summer!