

Paging

CS 537: Introduction to Operating Systems

Louis Oliphant

University of Wisconsin - Madison

Fall 2024

Administrivia

- P1 graded – see Piazza post
- P2 due today
- P3 out today
 - Large Project, due in 2 weeks
 - Go to Discussion!

Agenda

- Disadvantages of Segmentation Model
- Understand Paging Model
- Advantages/Disadvantages of Paging

Review: Segmentation

- Divides address space into logical segments (**code/static, stack, heap**)
- Have separate **base/bounds pairs for each segment**
- Segments grows and are placed independently
- Logical Addresses broken into segment ID and offset (for explicit approach)
- Extra protection bits and growth direction for segments

Review: Segmentation

```
./segmentation.py -c
```

```
address space size 1k
```

```
phys mem size 16k
```

```
Segment register information:
```

```
Segment 0 base (grows positive) : 0x1aea (decimal 6890)
```

```
Segment 0 limit : 472
```

```
Segment 1 base (grows negative) : 0x1254 (decimal 4692)
```

```
Segment 1 limit : 450
```

```
Virtual Address Trace
```

```
VA 0: 0x020b (decimal: 523) --> SEGMENTATION VIOLATION (SEG1)
```

```
VA 1: 0x019e (decimal: 414) --> VALID in SEG0: 0x1c88 (decimal: 7304)
```

```
VA 2: 0x0322 (decimal: 802) --> VALID in SEG1: 0x1176 (decimal: 4470)
```

```
VA 3: 0x0136 (decimal: 310) --> VALID in SEG0: 0x1c20 (decimal: 7200)
```

```
VA 4: 0x01e8 (decimal: 488) --> SEGMENTATION VIOLATION (SEG0)
```

For each virtual address, either write down the physical address it translates to OR write down that it is an out-of-bounds address (a segmentation violation). For this problem, you should assume a simple address space with two segments: the top bit of the virtual address can thus be used to check whether the virtual address is in segment 0 (topbit=0) or segment 1 (topbit=1). Note that the base/limit pairs given to you grow in different directions, depending on the segment, i.e., segment 0 grows in the positive direction, whereas segment 1 in the negative.


Quiz 5 : Segmentation Base/Bounds Translation

<https://tinyurl.com/cs537-fa24-q5>



Segmentation Disadvantages

- Each segment must be allocated contiguously
- May not have sufficient physical memory for large segments
- External Fragmentation (makes managing free memory hard!)



PAGING

Goal: Eliminate requirement that address space is contiguous


Eliminate external fragmentation

Grow segments as needed

Idea:

Divide address spaces and physical memory into fixed-sized pages


Size: 2^n , Example: 4KB



TRANSLATION OF PAGE ADDRESSES

How to translate logical address to physical address?

- High-order bits of address designate page number
- Low-order bits of address designate offset within page



No addition needed; just append bits correctly!

ADDRESS FORMAT

Given known page size, how many bits are needed in address to specify offset in page?

Page Size	Low Bits (offset)
16 bytes	
1 KB	
1 MB	
512 bytes	
4 KB	

ADDRESS FORMAT

Given number of bits in virtual address and bits for offset, how many bits for virtual page number?

Page Size	Low Bits(offset)	Virt Addr Total Bits	High Bits(vpn)
16 bytes	4	10	
1 KB	10	20	
1 MB	20	32	
512 bytes	9	16	
4 KB	12	32	

ADDRESS FORMAT

Given number of bits for vpn, how many virtual pages can there be in an address space?


Page Size	Low Bits (offset)	Virt Addr Bits	High Bits (vpn)	Virt Pages
16 bytes	4	10	6	
1 KB	10	20	10	
1 MB	20	32	12	
512 bytes	9	16	7	
4 KB	12	32	20	

VIRTUAL → PHYSICAL PAGE MAPPING

Number of bits in
virtual address

need not equal

number of bits in
physical address



How should OS translate VPN to PPN?

PAGETABLES


What is a good data structure ?

Simple solution: Linear page table aka *array*


VPN

0


2^n



PER-PROCESS PAGETABLE



FILL IN PAGETABLE



PageTable Size

Consider a **32-bit** address space with 4 KB pages. Assume each PTE is 4 bytes

How many bits do we need to represent the **offset** within a page?

How **many virtual pages** will we have in this case?

What will be the **overall size** of the page table?

WHERE ARE PAGETABLES STORED?

Implication: Store each page table in memory

Hardware finds page table base with register (e.g., CR3 on x86)

What happens on a context-switch?

Change contents of page table base register to newly scheduled process

Save old page table base register in PCB of descheduled process

OTHER PAGETABLE INFO

What other info is in pagetable entries besides translation?

- valid bit
- protection bits
- present bit (needed later)
- reference bit (needed later)
- dirty bit (needed later)

Pagetable entries are just bits stored in memory

- Agreement between HW and OS about interpretation

MEMORY ACCESSES WITH PAGING

14 bit addresses

0x0010: movl 0x1100, %edi

Assume PT is at phys addr 0x5000

Assume PTE's are 4 bytes

Assume 4KB pages

How many bits for offset? 12

Simplified view
of page table

2
0
80
99

Fetch instruction at logical addr 0x0010

Access page table to get ppn for vpn 0

Mem ref 1:

Learn vpn 0 is at ppn ____

Fetch instruction at _____ (Mem ref 2)

MEMORY ACCESSES WITH PAGING

14 bit addresses

0x0010: movl 0x1100, %edi

Assume PT is at phys addr 0x5000

Assume PTE's are 4 bytes

Assume 4KB pages

How many bits for offset? 12

Simplified view
of page table

2
0
80
99

Exec, load from logical addr 0x1100

Access page table to get ppn for vpn 1

Mem ref 3:

Learn vpn 1 is at ppn ____

Movl from ____ into reg (Mem ref 4)

MEMORY ACCESSES WITH PAGING

14 bit addresses

0x0010: movl 0x1100, %edi

Assume PT is at phys addr 0x5000

Assume PTE's are 4 bytes

Assume 4KB pages

How many bits for offset? 12

Simplified view
of page table

2
0
80
99

Fetch instruction at logical addr 0x0010

Access page table to get ppn for vpn 0

Mem ref 1: 0x5000

Learn vpn 0 is at ppn 2

Fetch instruction at 0x2010 (Mem ref 2)

Exec, load from logical addr 0x1100

Access page table to get ppn for vpn 1

Mem ref 3: 0x5004

Learn vpn 1 is at ppn 0

Movl from 0x0100 into reg (Mem ref 4)

PROS/CONS OF PAGING

No external fragmentation

Any page can be placed in any frame in physical memory

Fast to allocate and free

- Alloc: No searching for suitable free space
- Free: Doesn't have to coalesce with adjacent free space

Internal fragmentation

- Page size may not match process needs
- Wasted memory grows with larger pages

Additional memory reference to page table →

- Page table must be stored in memory
- MMU stores only base address of page table

Storage for page tables may be substantial

- Requires PTE for all pages in address space
- Entry needed even if page not allocated ?

SUMMARY: PAGE TRANSLATION STEPS

For each mem reference:

1. extract **VPN** (virt page num) from **VA** (virt addr)
2. calculate addr of **PTE** (page table entry)
3. read **PTE** from memory
4. extract **PFN** (page frame num)
5. build **PA** (phys addr)
6. read contents of **PA** from memory into register

Which steps are expensive?

- Notice that every memory reference requires 2 memory accesses (once to read the PTE and once to read the PA from memory)
- The page table must be stored in memory, increasing memory usage for every process

Next Two Lectures

- Using caching to reduce memory accesses (TLB)
- Shrinking the size of the page table (multi-level page tables)