

Concurrency: Threads

CS 537: Introduction to Operating Systems

Louis Oliphant

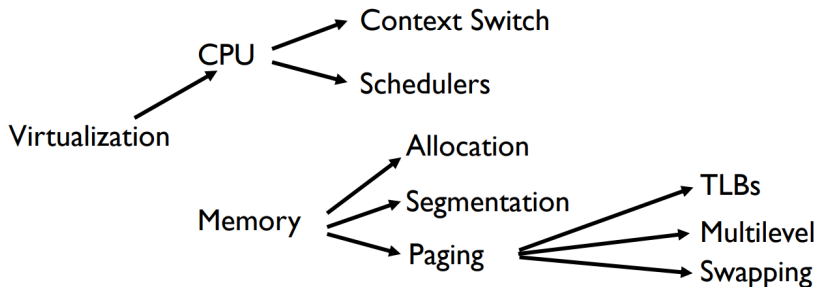
University of Wisconsin - Madison

Fall 2024

Administrivia

- Code reviews begin this Friday (look for an email from your TA)
- Midterm 1 Scheduled:
 - Regular Time: Oct 15th, 5:45-7:15
 - Last name A-K: Van Vleck B102
 - Last name L-Z: Ingraham B10
 - Unable to attend? Fill out [this form](#)
 - Alternate Time: Oct 16th, 5:45-7:15, Psych Bldg 121
 - McBurney Time: Oct 15th, 5:45-8:45pm, CS 1257
 - Bring **#2 Pencil**, one sheet of **8.5x11 (A4) notes**, and **UW Student ID**

Review: Virtualization



Correction: Bounds Check with Negative Growth

Segment	Base	Size (max 4K)	Grows Positive?
Code ₀₀	32K	2K	1
Heap ₀₁	34K	3K	1
Stack ₁₁	28K	2K	0

Figure 16.4: Segment Registers (With Negative-Growth Support)

Virtual address 15KB, which should map to physical address 27KB. Binary form is: 11 1100 0000 0000 (hex 0x3C00). Hardware uses top two bits (11) for segment, left with offset of 3KB. For the negative offset subtract maximum segment size from 3KB: **3KB minus 4KB which equals -1KB**. Add negative offset (-1KB) to the base (28KB) to get physical address: 27KB. **The bounds check can be calculated by ensuring the absolute value of the negative offset is less than or equal to the segment's current size (in this case, 2KB).**

Concurrency: Motivation

CPU Performance


Goal: Write applications that fully utilize many cores

Moore's Law: The number of transistors on microchips doubles every two years

Moore's law describes the empirical regularity that the number of transistors on integrated circuits doubles approximately every two years. This advancement is important for other aspects of technological progress in computing – such as processing speed or the price of computers.

Our World
in Data

Transistor count



Data source: Wikipedia ([wikipedia.org/wiki/Transistor_count](https://en.wikipedia.org/wiki/Transistor_count))

OurWorldinData.org – Research and data to make progress against the world's largest problems.

Licensed under CC-BY by the authors Hannah Ritchie and Max Roser.

Option 1: Communicating Processes

- Build Application using multiple processes
 - Example: Google Chrome (each tab is a process)
 - Communicate via `pipe()` or something similar
- Pros
 - Don't need new abstraction
 - Good for security
- Cons
 - Cumbersome programming
 - High communication overheads
 - Expensive context switch

Option 2: Threading

- New abstraction: **thread**
- Threads are like processes, except:
 - Multiple threads of same process **share an address space**
- Divide large task across several cooperative threads
- Communicate through **shared address space**

Common Concurrency Programming Models

- **Embarrassingly parallel**
 - e.g., ray tracing, database select, compiling separate files
- **Producer/Consumer**
 - Multiple producer threads create data (or work) that is handled by one of the multiple consumer threads
- **Pipeline**
 - Task is divided into series of subtasks, each processed on a different core
- **Defer work with background thread**
 - One thread performs non-critical work in the background (when CPU would be idle)


Thread vs. Process

Multiple threads share:

- Process ID (PID)
- Address space: Code (instructions), Most data (heap)
- Open file descriptors
- Current working directory
- User and group ID

Each thread has its own:

- Thread ID (TID)
- Set of registers, including PC and SP
- Stack for local vars and return address



```
void *mythread(void *arg) {
    printf("%s\n", (char *) arg);
    return NULL;
}

int main() {
    pthread_t p1, p2;
    printf("begin\n");
    Pthread_create(&p1, NULL, mythread, "A");
    Pthread_create(&p2, NULL, mythread, "B");
    // join waits for threads to finish
    Pthread_join(p1, NULL);
    Pthread_join(p2, NULL);
    printf("end\n");
    return 0;
}
```

Example Thread Trace 1 On Uniprocessor

main	Thread 1	Thread 2
start		
print "begin"		
create Thread 1		
create Thread 2		
wait for T1		
	print "A"	
wait for T2		
		print "B"
print "end"		

Example Thread Trace 2 On Uniprocessor

main	Thread 1	Thread 2
start		
print "begin"		
create Thread 1		
	print "A"	
create Thread 2		
		print "B"
wait for T1 (<i>returns immediately</i>)		
wait for T2 (<i>returns immediately</i>)		
print "end"		

Example Thread Trace 3 On Uniprocessor

main	Thread 1	Thread 2
start		
print "begin"		
create Thread 1		
create Thread 2		
		print "B"
wait for T1		
	print "A"	
wait for T2 (<i>returns immediately</i>)		
print "end"		

A Program's Behavior is *non-deterministic*.

Example Sharing Data

```
int max;
volatile int counter = 0; // shared global variable

void *mythread(void *arg) {
    char *letter = arg;
    int i; // stack (private per thread)
    printf("%s: begin [addr of i: %p]\n", letter, &i);
    for (i = 0; i < max; i++) {
        counter = counter + 1; // shared: only one
    }
    printf("%s: done\n", letter);
    return NULL;
}
```

```
int main(int argc, char *argv[]) {
    if (argc != 2) {
        fprintf(stderr, "usage: main-first <loopcount>\n");
        exit(1);
    }
    max = atoi(argv[1]);

    pthread_t p1, p2;
    printf("main: begin [counter = %d] [%x]\n", counter,
        (unsigned int) &counter);
    Pthread_create(&p1, NULL, mythread, "A");
    Pthread_create(&p2, NULL, mythread, "B");
    // join waits for the threads to finish
    Pthread_join(p1, NULL);
    Pthread_join(p2, NULL);
    printf("main: done\n [counter: %d]\n [should: %d]\n",
        counter, max*2);
    return 0;
}
```

```
prompt> ./threads 1000000
main: begin [counter = 0] [2df4a030]
A: begin [addr of i: 0x7fc6debfee3c]
B: begin [addr of i: 0x7fc6de3fde3c]
A: done
B: done
main: done
    [counter: 1013349]
    [should: 2000000]
```

Uncontrolled Scheduling – Race Condition

```
counter=counter+1;  // Critical Section
```

```
mov 0x8049a1c, %eax  
add $0x1, %eax  
mov %eax, 0x8049a1c
```

OS	Thread 1	Thread 2	(after instruction)		
			PC	eax	counter
	before critical section		100	0	50
	mov 8049a1c,%eax		105	50	50
	add \$0x1,%eax		108	51	50
	interrupt				
	save T1				
	restore T2		100	0	50
		mov 8049a1c,%eax	105	50	50
		add \$0x1,%eax	108	51	50
		mov %eax,8049a1c	113	51	51
	interrupt				
	save T2				
	restore T1		108	51	51
	mov %eax,8049a1c		113	51	51

What value is counter? Starting value = 50

Thread 1	Thread 2
<code>mov 0x8049a1c, %eax</code>	
<code>add \$0x1, %eax</code>	
<code>mov %eax, 0x8049a1c</code>	<code>mov 0x8049a1c, %eax</code>
	<code>add \$0x1, %eax</code>
	<code>mov %eax, 0x8049a1c</code>

Thread 1	Thread 2
	<code>mov 0x8049a1c, %eax</code>
	<code>add \$0x1, %eax</code>
	<code>mov %eax, 0x8049a1c</code>
<code>mov 0x8049a1c, %eax</code>	
<code>add \$0x1, %eax</code>	
<code>mov %eax, 0x8049a1c</code>	

Thread 1	Thread 2
	<code>mov 0x8049a1c, %eax</code>
<code>mov 0x8049a1c, %eax</code>	
	<code>add \$0x1, %eax</code>
<code>add \$0x1, %eax</code>	
	<code>mov %eax, 0x8049a1c</code>
<code>mov %eax, 0x8049a1c</code>	

What value is counter? Starting value = 50

Thread 1	Thread 2	
mov 0x8049a1c, %eax		
add \$0x1, %eax		
	mov 0x8049a1c, %eax	
mov %eax, 0x8049a1c	add \$0x1, %eax	
	mov %eax, 0x8049a1c	counter = 51

Thread 1	Thread 2	
	mov 0x8049a1c, %eax	
	add \$0x1, %eax	
	mov %eax, 0x8049a1c	
mov 0x8049a1c, %eax		
add \$0x1, %eax		
mov %eax, 0x8049a1c		counter = 52

Thread 1	Thread 2	
	mov 0x8049a1c, %eax	
mov 0x8049a1c, %eax		
	add \$0x1, %eax	
add \$0x1, %eax		
mov %eax, 0x8049a1c	mov %eax, 0x8049a1c	counter = 51

Non-Determinism

Concurrency leads to **non-deterministic** results

- Different results even with same inputs
- **Race Condition** – results depend upon the scheduling order

Whether bug manifests depends on CPU scheduling!

What We Want

Want 3 instructions to execute as an uninterruptable group

```
mov 0x8049a1c, %eax  
add $0x1, %eax  
mov %eax, 0x8049a1c
```

Want them to be **Atomic** – “as a unit” or “all or nothing”. The three instructions should all run together (or not at all).

Synchronization Primitives

- Hardware support helps to build **Synchronization primitives**:
 - **Lock**
 - **Condition Variable**
 - **Semaphore**
 - **Barrier**
- Used to create atomicity for critical sections
- Also used to make one thread wait for another thread to complete some action before continuing

Why in OS Class?

- OS is the first concurrent program
- Page tables, process lists, file system structures, and most kernel data must be accessed using proper synchronization primitives.

Thread Creation

```
#include <pthread.h>

typedef struct __myarg_t {
    int a;
    int b;
} myarg_t;

void *mythread(void *arg) {
    myarg_t *m = (myarg_t *) arg;
    printf("%d %d\n", m->a, m->b);
    return NULL;
}

int main(int argc, char *argv[]) {
    pthread_t p;
    myarg_t args;
    args.a = 10;
    args.b = 20;
    rc = pthread_create(&p, NULL, mythread, &args); //success returns 0
}
```

Thread Joining (and returning values)

```
typedef struct __myret_t {
    int x;
    int y;
} myret_t;

void *mythread(void *arg) {
    ...
    myret_t *r = Malloc(sizeof(myret_t));
    r->x = 1;
    r->y = 2;
    return (void *) r;
}

int main(int argc, char *argv[]) {

    myret_t *m;
    ...
    rc=pthread_create(...);
    rc=pthread_join(p, (void **) &m); //success returns 0
    printf("returned %d %d\n", m->x, m->y);
    free(m);
}
```