

# Concurrency: Locks

## CS 537: Introduction to Operating Systems

Louis Oliphant

University of Wisconsin - Madison

Fall 2024

# Administrivia

- CS Midterm Evaluation Fall 2024 Survey (Check Your Emails)
- Code reviews this week
- Midterm 1 Scheduled TONIGHT:
  - Regular Time: Oct 15th, 5:45-7:15 - Last name A-K: Van Vleck B102 - Last name L-Z: Ingraham B10
  - McBurney Time: Oct 15th, 5:45-8:45pm, CS 1257
  - Alternate Time: Oct 16th, 5:45-7:15, Psych Bldg 121
  - Bring **#2 Pencil**, one sheet of **8.5x11 (A4) notes**, and **UW Student ID**
  - Turn in exam booklet, sheet of notes, and scantron.

## Review: Threads

A **thread** is similar to a process in that it is a unit of execution and can be scheduled.

The main difference is threads **share virtual address space** (code and heap data) but have different **registers and stack**.

Understand the **race condition** between threads accessing shared data.

Remember how to:

- Create a thread
- Wait for a thread to finish executing
- Pass arguments to a thread
- Get return values after a thread has finished

## Review Threads: looping-race-nolock.s

```
# assumes %bx has loop count in it
.main
.top
# critical section
mov 2000, %ax # get 'value' at address 2000
add $1, %ax   # increment it
mov %ax, 2000 # store it back

# see if we're still looping
sub $1, %bx
test $0, %bx
jgt .top

halt
```

# Review: x86.py

```
./x86.py -p looping-race-nolock.s --argv='bx=2' --threads=1 --memtrace=2000 -c
```

```
2000          Thread 0
0
0 1000 mov 2000, %ax
0 1001 add $1, %ax
1 1002 mov %ax, 2000
1 1003 sub $1, %bx
1 1004 test $0, %bx
1 1005 jgt .top
1 1000 mov 2000, %ax
1 1001 add $1, %ax
2 1002 mov %ax, 2000
2 1003 sub $1, %bx
2 1004 test $0, %bx
2 1005 jgt .top
2 1006 halt
```

```
./x86.py -p looping-race-nolock.s --argv='bx=2' --threads=2 --memtrace=2000 -c
```

2000	Thread 0	Thread 1
0		
0	1000 mov 2000, %ax	
0	1001 add \$1, %ax	
1	1002 mov %ax, 2000	
1	1003 sub \$1, %bx	
1	1004 test \$0, %bx	
1	1005 jgt .top	
1	1000 mov 2000, %ax	
1	1001 add \$1, %ax	
2	1002 mov %ax, 2000	
2	1003 sub \$1, %bx	
2	1004 test \$0, %bx	
2	1005 jgt .top	
2	1006 halt	
2	----- Halt;Switch -----	----- Halt;Switch -----
2		1000 mov 2000, %ax
2		1001 add \$1, %ax
3		1002 mov %ax, 2000
3		1003 sub \$1, %bx
3		1004 test \$0, %bx
3		1005 jgt .top
3		1000 mov 2000, %ax
3		1001 add \$1, %ax
4		1002 mov %ax, 2000
4		1003 sub \$1, %bx
4		1004 test \$0, %bx
4		1005 jgt .top
4		1006 halt

```
./x86.py -p looping-race-nolock.s -a 'bx=2' -t 2 -M 2000 --regtrace=ax --interrupt=6 --randints -s 3 -c
```

2000	ax	Thread 0	Thread 1	2000	ax	Thread 0	Thread 1
0	0			2	2		1002 mov %ax, 2000
0	0	1000 mov 2000, %ax		2	2	--- Interrupt ----	---- Interrupt ----
0	1	1001 add \$1, %ax		2	2	1003 sub \$1, %bx	
0	0	--- Interrupt ----	---- Interrupt ----	2	2	--- Interrupt ----	---- Interrupt ----
0	0		1000 mov 2000, %ax	2	2		1003 sub \$1, %bx
0	1		1001 add \$1, %ax	2	2		1004 test \$0, %bx
1	1		1002 mov %ax, 2000	2	2		1005 jgt .top
1	1		1003 sub \$1, %bx	2	2		1006 halt
1	1	--- Interrupt ----	---- Interrupt ----	2	2	-- Halt;Switch ---	--- Halt;Switch ---
1	1	1002 mov %ax, 2000		2	2	1004 test \$0, %bx	
1	1	1003 sub \$1, %bx		2	2	1005 jgt .top	
1	1	1004 test \$0, %bx		2	2	--- Interrupt ----	---- Interrupt ----
1	1	--- Interrupt ----	---- Interrupt ----	2	2	1006 halt	
1	1		1004 test \$0, %bx				
1	1		1005 jgt .top				
1	1		1000 mov 2000, %ax				
1	2		1001 add \$1, %ax				
1	1	--- Interrupt ----	---- Interrupt ----				
1	1	1005 jgt .top					
1	1	1000 mov 2000, %ax					
1	2	1001 add \$1, %ax					
2	2	1002 mov %ax, 2000					
2	2	--- Interrupt ----	---- Interrupt ----				

## Quiz 9: Intro to Threads

<https://tinyurl.com/cs537-fa24-q9>





# Locks (Programmer's Perspective)

```
#include <pthread.h>

pthread_mutex_t lock;
pthread_mutex_lock(&lock);
x=x+1; //or whatever your critical section is
pthread_mutex_unlock(&lock);
```

Use different locks to protect different variables / data structures

Rest of lecture is to understand how locks are built (hardware and OS support)

Next lecture on using locks with different data structures

# Lock Implementation Goals

- ① **Safety:** *mutual exclusion*, only one thread in critical section at a time
- ② **Liveness:**
  - *deadlock free*: if two simultaneous requests, must allow one to proceed
  - *starvation free*: must eventually allow each waiting thread to enter
  - *fairness*: each thread waits for same amount of time
- ③ **Performance** – minimize CPU usage for lock/unlock
  - case 1 – no contention
  - case 2 – multiple threads contending, single CPU
  - case 3 – multiple threads contending, multiple CPUs

# Earliest Solution: Disable Interrupts

```
void lock() {  
    DisableInterrupts();  
}  
  
void unlock() {  
    EnableInterrupts();  
}
```

On single CPU, thread assured no other thread will interfere (including OS)

**This approach used sparingly by OS itself**

## Disadvantages

- User program has control of CPU, could `lock()` and run forever
- Doesn't work on multiprocessor systems
- Can lead to lost interrupts (imagine OS not being notified of I/O completion)
- Very inefficient

# Failed Attempt: Using Loads/Stores

Why doesn't this work?

```
typedef struct __lock_t {int flag; } lock_t;

void init(lock_t *mutex) {
    //0 -> lock is available, 1 -> held
    mutex->flag = 0;
}

void lock(lock_t *mutex) {
    while (mutex->flag == 1) { // TEST the flag
    }
    mutex->flag = 1;
}

void unlock(lock_t *mutex) {
    mutex->flag = 0;
}
```

## Failed Attempt: Reason

### Thread 1

```
call lock ()
while (flag == 1)
interrupt: switch to Thread 2
```

```
flag = 1; // set flag to 1 (too!)
```

### Thread 2

```
call lock ()
while (flag == 1)
flag = 1;
interrupt: switch to Thread 1
```

## No Mutual Exclusion!

### Wasteful Spin-waiting

# Hardware Support: Atomic Exchange (book calls TestAndSet())

On x86 `xchg (dst), src` atomically swaps the contents of memory at `dst` with register `src`

## Happens Atomically:

```
int xchg(int *addr, int new) {  
    int old = *addr;  
    *addr = new;  
    return old;  
}
```

```
int TestAndSet(int *old_ptr, int new) {  
    int old = *old_ptr;  
    *old_ptr = new;  
    return old;  
}
```

# Lock Implementation with Xchg

```
typedef struct __lock_t {
    int flag;
} lock_t;

void init(lock_t *lock) {
    lock->flag = 0;
}

void lock(lock_t *lock) {
    while(Xchg(&lock->flag, 1) == 1) {
        // spin-wait (do nothing)
    }
}

void unlock(lock_t *lock) {
    lock->flag = 0;
}
```

## Other Atomic HW Instructions

```
int CompareAndSwap(int *addr, int expected, int new) {  
    int actual = *addr;  
    if (actual == expected) {  
        *addr = new;  
    }  
    return actual;  
}
```

```
void lock(lock_t *lock) {  
    while (CompareAndSwap(&lock->flag, 0, 1) != 0) {  
        //spin  
    }  
}
```



# A Pair of Atomic HW Instructions


```
int LoadLinked(int *ptr) {
    return *ptr;
}

int StoreConditional(int *ptr, int value) {
    if (no one has updated *ptr since the LoadLinked to this address) {
        *ptr = value;
        return 1; //success
    } else {
        return 0; //failed to update
    }
}

void lock(lock_t *lock) {
    while (1) {
        while (LoadLinked(&lock->flag) == 1) {
            // spin until it's zero
        }
        if (StoreConditional(&lock->flag, 1) == 1) {
            return; //if set-it-to-1 was a success: all done
                //otherwise: try it all over again
        }
    }
}

void unlock(lock_t *lock) {
    lock->flag = 0;
}
```

# Basic Spinlocks Are Unfair



**Scheduler is unaware of locks/unlocks!**

# Fairness: Ticket Locks – Based on Atomic HW Instruction

```
int FetchAndAdd(int *ptr) {  
    int old = *ptr;  
    *ptr = old + 1;  
    return old;  
}
```

```
typedef struct __lock_t {  
    int ticket; //thread's ticket number  
    int turn; //whose turn it is  
} lock_t;  
  
void lock_init(lock_t *lock) {  
    lock->ticket = 0;  
    lock->turn = 0;  
}  
  
void lock(lock_t *lock) {  
    //first, reserve this thread's turn  
    int myturn = FetchAndAdd(&lock->ticket);  
    while (lock->turn != myturn)  
        ; //spin until thread's turn  
}  
  
void unlock(lock_t *lock) {  
    lock->turn = lock->turn+1;  
}
```

# Ticket Lock Example

A lock():

B lock():

C lock():


A unlock():

A lock():

B unlock():

C unlock():

A unlock():



# Spinlock Performance


Fast when...

- many CPUs
- locks held a short time
- advantage: avoid context switch

Slow when...

- one CPU
- locks held a long time
- disadvantage: spinning is wasteful

## CPU Scheduler is Ignorant of Spinlocks



CPU scheduler may run **B,C,D** instead of **A**  
even though **B,C,D** are waiting for **A**


# Ticket Lock With Yield (OS Call)

Instead of spinning, give up CPU with special yield() instruction

```
typedef struct __lock_t {
    int ticket; //thread's ticket number
    int turn; //whose turn it is
}
void lock_init(lock_t *lock) {
    lock->ticket = 0;
    lock->turn = 0;
}
void lock(lock_t *lock) {
    int myturn = FetchAndAdd(&lock->ticket); // Reserve turn
    while (lock->turn != myturn)
        yield(); //give up rest of time-slice
}

void unlock(lock_t *lock) {
    lock->turn = lock->turn+1;
}
```

# Yield Instead of Spin





# Spinlock Performance

Waste of CPU cycles?

- Without yield: threads \* **time\_slice**
- With yield: threads \* **context\_switch\_time**

Even with yield, spinning is slow with high thread contention

Next improvement: put thread on waiting queue and block instead of spinning

New OS call:

- `park()` – put calling thread to sleep
- `unpark()` – wake a particular thread

# Lock Implementation (Two-Phase Locks)

```
typedef struct __lock_t {
    int guard; // guards flag and q
    int flag; // 1 if lock is acquired
    queue_t *q;
}

void lock_init(lock_t *m) {
    m->flag = 0;
    m->guard = 0;
    queue_init(m->q);
}

void lock(lock_t *m) {
    while (Xchg(&m->guard, 1) == 1)
        ; //acquire guard lock by spinning
    if (m->flag == 0) {
        m->flag = 1; //lock is acquired
        m->guard = 0;
    } else {
        queue_add(m->q, getpid());
        m->guard = 0;
        park();
    }
}
```

```
void unlock(lock_t *m) {
    while (Xchg(&m->guard, 1) == 1)
        ; //acquire guard lock by spinning
    if (queue_empty(m->q))
        m->flag = 0; //let go of lock, no one wants it
    else
        unpark(queue_remove(m->q)); //hold for next
    m->guard = 0;
}
```

What would happen if release of guard came *after* the park()?

Think about possible **wakeup/waiting race condition** just before the call to park()

Add setpark() OS call to indicate *about* to park(). Add call to setpark() just before releasing guard:

```
queue_add(m->q, getpid());
setpark();
m->guard = 0;
park();
```

# Lock Implementation: fixing race with setpark

```
void lock(lock_t *m) {
    while (Xchg(&m->guard, 1) == 1)
        ; //acquire guard lock by spinning
    if (m->flag == 0) {
        m->flag = 1; //lock is acquired
        m->guard = 0;
    } else {
        queue_add(m->q, getpid());
        m->guard = 0;
        park();
    }
}

void unlock(lock_t *m) {
    while (Xchg(&m->guard, 1) == 1)
        ; //acquire guard lock by spinning
    if (queue_empty(m->q))
        m->flag = 0; //let go of lock, no one wants it
    else
        unpark(queue_remove(m->q)); //hold for next
    m->guard = 0;
}
```

Just before `park()`, this code has a **wakeup/waiting race condition**

Add `setpark()` OS call to indicate *about* to `park()`:

```
queue_add(m->q, getpid());
setpark();
m->guard = 0;
park();
```

After `setpark()`, if `unpark()` is called first then subsequent `park()` returns immediately.

# Spin-Waiting vs. Blocking

Each approach is better under different circumstances:

- Uniprocessor
  - Waiting process is scheduled  $\rightarrow$  Process holding lock is not
  - Waiting process should always relinquish processor
  - Associate queue of waiters with each lock (as in previous implementation)
- Multiprocessor
  - Waiting process is scheduled  $\rightarrow$  process holding lock might be
  - spin or block depends on  $t$  time before lock released vs context-switch cost  $C$ :
    - Lock released quickly ( $t \ll C$ )  $\rightarrow$  **Spin-wait**
    - Lock released slowly ( $t \geq C$ )  $\rightarrow$  **Block**