

# Concurrency: Common Problems

## CS 537: Introduction to Operating Systems

Louis Oliphant

University of Wisconsin - Madison

Fall 2024

# Administrivia

- Project 3 Graded
- Project 4 due Tue Nov 5th @ 11:59pm
- Exam 2, Thu, Nov 7th 5:45-7:15pm
  - Conflict? – fill out the [exam conflict form](#)
  - Same format as Exam 1
  - Exam 2 Review next lecture – Come with questions

## Review: Semaphores

A **semaphore** is an object with an integer value that must be initialized and can be manipulated with two routines:

```
#include <semaphore.h>
sem_t s;
sem_init(&s, 0, 1); //initializes to 1 (3rd arg)
```

```
int sem_wait(sem_t *s) {
    decrement the value of semaphore s by one
    wait if value of semaphore s is negative
}
```

```
int sem_post(sem_t *s) {
    increment the value of semaphore s by one
    if there are threads waiting, wake one
}
```

## Quiz: Semaphores


<https://tinyurl.com/cs537-fa24-q13>



# Concurrency Problems Agenda

- Non-Deadlock Bugs
  - Atomicity Violation
  - Order Violation
- Deadlock Bugs
  - Why they occur
  - How to prevent them

# Concurrency Study



Lu *et al.* [ASPLOS 2008]:

- For four major projects, search for concurrency bugs among >500K bug reports
- Analyze small sample to identify common types of concurrency bugs

# Atomicity Violation

## Thread 1:

```
if (thd->proc_info) {  
    ...  
    fputs(thd->proc_info,...);  
    ...  
}
```

## Thread 2:

```
thd->proc_info = NULL;
```

# Fix Atomicity Violations with Locks

## Thread 1:

```
pthread_mutex_lock(&lock);  
if (thd->proc_info) {  
    ...  
    fputs(thd->proc_info,...);  
    ...  
}  
pthread_mutex_unlock(&lock);
```

## Thread 2:

```
pthread_mutex_lock(&lock);  
thd->proc_info = NULL;  
pthread_mutex_unlock(&lock);
```



# Order Violation

## Thread 1:

```
void init() {  
    ...  
    mThread =  
        PR_CreateThread(mMain,...);  
    ...  
}
```

## Thread 2:

```
void mMain(...) {  
    ...  
    mState = mThread->State;  
    ...  
}
```

# Fix Order Violations with Condition Variables

## Thread 1:


```
void init() {  
    ...  
    mThread =  
        PR_CreateThread(mMain,...);  
    pthread_mutex_lock(&mtLock);  
    mtInit = 1;  
    pthread_cond_signal(&mtCond);  
    pthread_mutex_unlock(&mtLock);  
    ...  
}
```

## Thread 2:

```
void mMain(...) {  
    ...  
    mutex_lock(&mtLock);  
    while (mtInit == 0)  
        Cond_wait(&mtCond, &mtLock);  
    Mutex_unlock(&mtLock);  
    mState = mThread->State;  
    ...  
}
```

# Why Deadlocks Occur

|                            |                            |
|----------------------------|----------------------------|
| Thread 1:                  | Thread 2:                  |
| <code>lock(&amp;A);</code> | <code>lock(&amp;B);</code> |
| <code>lock(&amp;B);</code> | <code>lock(&amp;A);</code> |




No progress can be made because two or more threads are waiting for the other to take some action and thus neither ever does (**Circular Dependency**).

## Fix by Removing Circular Dependencies

Have an order that locks are obtained:

|                            |                            |
|----------------------------|----------------------------|
| Thread 1:                  | Thread 2:                  |
| <code>lock(&amp;A);</code> | <code>lock(&amp;A);</code> |
| <code>lock(&amp;B);</code> | <code>lock(&amp;B);</code> |



# Conditions for Deadlock

- ① **Mutual Exclusion** – Threads claim exclusive control of resources that they require (e.g. a thread grabs a lock)
- ② **Hold-and-wait** – Threads hold resources allocated to them while waiting for additional resources
- ③ **No preemption** – Resources cannot be forcibly removed from threads
- ④ **Circular wait** – Circular chain of threads hold resources that other threads need

Remove *any one* of these criteria and deadlock cannot occur.

## Problem: encapsulation

Solving deadlocks becomes trickier with **encapsulated** code, e.g. Vector class in Java:

Thread 1

```
v1.addAll(v2);
```

Thread 2

```
v2.addAll(v1);
```

Need to know that `v1.addAll(v2)` acquires locks on `v1` and `v2`.

Can't control how Vector acquires locks.

## Prevention Technique – Lock ordering (condition 4)

Create a **total** or **partial** lock ordering

- can order by lock address in memory
- requires programmer discipline

Problem: might not know what locks are needed upfront

## Prevention Technique – Hold-and-wait (condition 2)

- Acquire all locks at once:

```
pthread_mutex_lock(prevention);  
pthread_mutex_lock(L1);  
pthread_mutex_lock(L2);  
...  
pthread_mutex_unlock(prevention);
```

- Can be problematic:
  - Encapsulation (must know what locks are required for each function call and get them)
  - Decreases concurrency since all locks must be acquired at once



## Prevention Technique – No Preemption (condition 3)

- Stop holding onto lock if you can't acquire the other needed locks:

```
top:
    pthread_mutex_lock(L1);
    if (pthread_mutex_trylock(L2) != 0) {
        pthread_mutex_unlock(L1);
        goto top;
    }
```

- New Problem: **Livelock** – two threads can repeatedly attempt this sequence and repeatedly fail to acquire both locks.
- Encapsulation still a problem (if a lock acquisition is buried in some routine, difficult to jump back).

# Final Prevention Technique – Avoid Mutual Exclusion (condition 1)

- Avoid needing mutual exclusion by using thread safe, **lock-free** data structures. These use the hardware atomic instructions.

```
// hardware guarantees atomicity
int CompareAndSwap(int *a, int e, int new) {
    if (*a == e) {
        *a = new;
        return 1; //success
    }
    return 0; //failure
}
```

Insert into a List:

```
void insert(int value) {
    node_t *n = malloc(sizeof(node_t));
    n->value = value;
    do {
        n->next = head;
    } while (CompareAndSwap(&head,
                           n->next, n) == 0);
}
```

# Other Strategies

## Deadlock Avoidance

A smart scheduler that is aware of which threads require which locks can schedule threads such that deadlock cannot occur.

## Deadlock Recovery

Allow deadlocks to occur (hopefully occasionally), have process to detect a deadlock, and then take some action to fix it.

# Concurrency Summary

- **Hardware support**
- **Threads** and shared memory
- **Locks** and protection surrounding critical code sections
  - Use of Locks to create **thread-safe data structures**
  - Use of hardware primitives for thread safety
- **Condition Variables** controlling thread execution / sleeping on some program state.
- **Semaphores** are flexible primitives that can replace locks and condition variables
- Use concurrency primitives to prevent common concurrency problems like **deadlock**, **starvation**, guarantee **atomicity** and thread order.