

# Persistence: Journaling

## CS 537: Introduction to Operating Systems

Louis Oliphant

University of Wisconsin - Madison

Fall 2024

# Administrivia

- Project 5 due Nov 19 @ 11:59pm
- Exam 2 [regrade form](#)

# Review File System Implementation

- Data Structures
  - Superblock, inode and data bitmap, inode table, data blocks, directory entries
- Access Methods
  - How does a call like `open()`, `read()`, or `write()` get mapped onto the data structures of the disk?

# Review Fast File System

- Treat the disk like it's a disk
  - Divide disk into **groups**
  - Each group gets superblock, block bitmap, inode bitmap, inode table, and data blocks
- *Keep related stuff together, keep unrelated stuff far apart*
  - Place directories in group with low number of directories but high number of free inodes
  - Place files (both data and inode) in group with directory
- Large File Exception

## Quiz 17 Fast File System

<https://tinyurl.com/cs537-fa24-q17>



## This lecture: crash consistency


A crash could interrupt the system between any two writes.  
How to **update the file system safely** despite this?

# fsck and Journaling

- Crash Consistency Problem
- Solution 1: fsck
- Solution 2: Journaling
  - Data Journaling, Recovery, Metadata Journaling
- Solution 3: Other Approaches

# Consistency in File System

The file system consists of several data structures which need to be **consistent with each other**.





# Consistency in File System

Consider the case of appending a single data block to an existing file. The following changes must occur:


- Update block bitmap to acquire a free block
- Update inode to  $I[v2]$  and point to new data block
- Write data  $Db$  to new data block

If a crash happens after one or two of these writes have taken place, but not all three, the file system could be left in an **inconsistent state**.

**Before:**



**After:**



## 3 Crash Scenarios

- **Just the data block (Db) is written to disk**
  - Since the inode and data bitmap are not updated, the block is not attached to the file and could still be allocated.
  - FS is consistent but does not contain the new data in Db.
- **Just the updated inode (I[v2]) is written to disk**
  - Since inode is updated but data was not written to Db, garbage is attached to the end of the file
  - Since data bitmap was not modified, block could be allocated to a different file (**file system inconsistency**)
- **Just the updated bitmap is written to disk**
  - data block has not been attached to the file, resulting in a **space leak**


We want file system updates to happen **atomically** from one consistent state to another.

## Solution #1: The File System Checker: `fsck`

Let inconsistencies happen, then fix by running `fsck`:


- First check **superblock**, that it looks reasonable.
- **Free blocks**: scan the inodes, direct blocks, indirect blocks, etc., to build a list of allocated and free data blocks. Compare this list with the data bitmap. Do the same for the inode bitmap compared to the inode table.
- **Inode state**: check for corruption, e.g. valid type field.
- **Inode links**: verify the link count
- **Duplicates**: Check for duplicate pointers to same data blocks
- **Bad Blocks**: Check for pointers outside range of valid data blocks
- **Directory Checks**: Integrity check (each directory contains a `.` and `..` pointer to proper values, inodes exist, each directory linked to once).

## Free Blocks Example




Block is unallocated but in inode

## Link Count Example



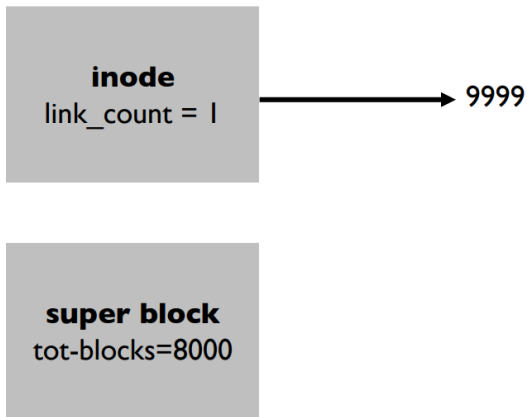
Link count is wrong

## Duplicate Pointers



Block is in two inodes

## Bad Pointer




Pointer is to invalid block number

# Problems with fsck

- Not always obvious how to fix
- Don't know “correct” state, just consistent state



# Fundamental problem with fsck: too slow



Checking a 600GB disk takes ~70 minutes


ffsck: fast file system checker

Ao Ma, Chris Dragg, Andrea C. Arpaci-Dusseau, and Remzi H. Arpaci-Dusseau

## Solution #2: Journaling (or Write-Ahead Logging)


Before updating the disk, first write down a little note (in the **journal** on disk) describing what you are about to do.

On a crash during update, go back and look at the journal.



Adds a bit of work during update but greatly reduces work required during recovery

# Data Journaling



- 1 **Journal write** – Write out TxB, I[v2], B[v2], and Db, wait for these to complete
- 2 **Journal commit** – Write out TxE, wait for this to complete, transaction is committed
- 3 **Checkpoint** – Write the contents of the update to their final on-disk locations

The TxB block contains information about the pending update (e.g. final addresses for blocks and a **transaction identifier**)

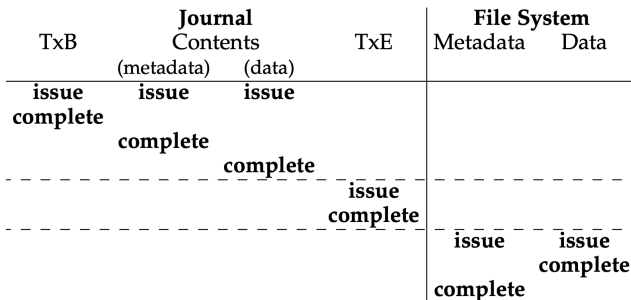
The disk guarantees any 512-byte write is atomic – the TxE is a single 512-bytes.

# Recovery

Crashes can occur at any time:

- If occurs before the commit, the update is skipped
- If occurs after the commit but before the checkpoint finishes:
  - On boot, scan the log and look for committed transactions and **replay** them in order: called **redo logging**
  - In the worse case, a transaction might be performed again

# Data journaling timeline



## Data Journaling

- Note the **write issues** which can occur simultaneously
- Completion time is determined by the I/O subsystem, which may reorder writes to improve performance
- Horizontal dashed lines representing barriers waiting for completion of writes are enforced by FS for protocol correctness

## Batching Log Updates


- This journaling protocol adds a lot of extra disk traffic
- Some FS do not commit each update one at a time, rather buffer all updates into a **global transaction**

## Making Log Finite

Journaling file systems treat the log as a **circular log** adds a 4th step to protocol:

- ④ **Free:** some time later, mark the transaction free in the journal by updating the journal superblock

**Journal Superblock** keeps track of portion of journal with non-checkpointed transactions.



# Metadata Journaling

Data journaling has high cost: each data block is written twice


**Ordered journaling** or **metadata journaling** only puts metadata in journal

More complicated to get right

# Metadata Journaling Protocol

By forcing data write first, a file system can guarantee that a pointer will never point to garbage – commonly used technique.

- 1 **Data write:** Write data to final location; wait for completion (optional)
- 2 **Journal metadata write:** Write the begin block and metadata to log; wait for completion
- 3 **Journal commit:** Write the transaction commit block (containing TxE)
- 4 **Checkpoint metadata:** Write the contents of the metadata to final locations
- 5 **Free:** Later, mark the transaction free in journal superblock





## Tricky Case: Block Reuse

- User creates a file in directory `foo`. FS journal has a write to `foo`'s inode and its block 1000.
- User deletes `foo`, FS frees block 1000.
- User creates a new file which FS gives block 1000. FS directly writes to block 1000.

On crash, replay write to block 1000 with old `foo` contents, but these are now owned by a file

Linux's ext FS has **revoke** entries in log. Replaying first checks for these entries and doesn't replay entries that are revoked

# Metadata journaling

## Metadata Journaling

TxB	Journal Contents (metadata)	TxE	File System	
			Metadata	Data
issue	issue			issue
				complete
complete	complete			
		issue		
		complete		
			issue	
			complete	

- Data is written once
- ext4 doesn't wait for data to be written - tricky to get right

## Solution #3: Other Approaches / Ideas

- **Copy-On-Write (COW)**

- Never overwrite in place: write new data, then point to it
- Basis of **Log File System** (lecture next time), and ZFS and btrfs

- **Soft Updates**

- Carefully order all writes to the FS
- Requires intricate knowledge of each FS data structure

- **Backpointer-based Consistency (BBC)**

- Developed here at UW
- Every block contains a back-pointer, so data blocks point to the inode they belong to.
- Consistency can be checked between the forward pointers in the inode with the backpointers

# Summary

- Crash consistency is a key problem in file systems
- Journaling is a key technique to make crash consistency easy