

# Advanced Topics: Distributed Systems Intro

## CS 537: Introduction to Operating Systems

Louis Oliphant

University of Wisconsin - Madison

Fall 2024

# Administrivia

- Project 6 - due Dec 6th
- Final Exam: 12/19 10:05-12:05
  - Humanities room 3650
  - Microbial Sciences room 1220
  - McBurney Accommodations, CS room 1325

# Review Multiprocessor Scheduling

- Multiprocessor Architecture
  - cache and cache coherency
- Multiprocessor Complications
  - cache affinity
  - synchronization
- SQMS and MQMS
- CFS Scheduler
- ULE Scheduler

## Quiz 22 SSDs

<https://tinyurl.com/cs537-fa24-quiz22>



# Distributed Systems

## Building Distributed Systems That Work When Components Fail

- System objectives of performance, security, **communication**
- Unreliable Communication Layers (**UDP**)
  - **checksum**
- Reliable Communication Layers (**TCP**)
  - acknowledgement, timeout/retry
  - sequence counter
- Communication Abstractions
  - Distributed Shared Memory (**DSM**)
  - Remote Procedure Call (**RPC**)
    - Stub Generator
    - Run-Time Library
    - Other Issues: fragmentation/reassembly, byte ordering, synchronicity

# WHAT IS A DISTRIBUTED SYSTEM?

*A distributed system is one where a machine I've never heard of can cause my program to fail.*

— [Leslie Lamport](#)

Definition: More than one machine working together to solve a problem

Examples:

- client/server: web server and web client
- cluster: page rank computation

# WHY GO DISTRIBUTED?

More computing power

More storage capacity

Fault tolerance

Data sharing

# NEW CHALLENGES

System failure: need to worry about **partial** failure

Communication failure: links unreliable

- bit errors
- packet loss
- node/link failure



# COMMUNICATION OVERVIEW

Raw messages: UDP

Reliable messages: TCP

Remote procedure call: RPC

# RAW MESSAGES: UDP

UDP : User Datagram Protocol

API:

- reads and writes over socket file descriptors
- messages sent from/to ports to target a process on machine

Provide minimal reliability features:

- messages may be lost
- messages may be reordered
- messages may be duplicated
- only protection: checksums to ensure data not corrupted

# RAW MESSAGES: UDP

## Advantages

- Lightweight
- Some applications make better reliability decisions themselves (e.g., video conferencing programs)

## Disadvantages

- More difficult to write applications correctly


# RELIABLE MESSAGES: LAYERING STRATEGY

TCP: Transmission Control Protocol

Using software to build

reliable logical connections over unreliable physical connections

# TECHNIQUE #1: ACK



Ack: Sender knows message was received  
What to do about message loss?

## TECHNIQUE #2: TIMEOUT

Sender

[send message]

[start timer]

... waiting for ack ...

[timer goes off]

[send message]

[recv ack]

Receiver



[recv message]

[send ack]

# TIMEOUT

How long to wait?


Too long?

- System feels unresponsive

Too short?

- Messages needlessly re-sent
- Messages may have been dropped due to overloaded server. Resending makes overload worse!

# LOST ACK PROBLEM





# SEQUENCE NUMBERS

## Sequence numbers

- senders gives each message an increasing unique seq number
- receiver knows it has seen all messages before N

Suppose message K is received.

- if  $K \leq N$ , Msg K is already delivered, ignore it
- if  $K = N + 1$ , first time seeing this message
- if  $K > N + 1$  ?

# TCP

TCP: Transmission Control Protocol

Most popular protocol based on seq nums

Buffers messages so arrive in order

Timeouts are adaptive

# COMMUNICATIONS OVERVIEW

Raw messages: UDP

Reliable messages: TCP

Remote procedure call: RPC

# RPC

## **R**emote **P**rocedure **C**all

What could be easier than calling a function?

**Approach:** create wrappers so calling a function on another machine feels just like calling a local function!

# RPC

## Machine A

```
int main(...) {  
    int x = foo("hello");  
}  
  
int foo(char *msg) {  
    send msg to B  
    recv msg from B  
}
```

## Machine B

```
int foo(char *msg) {  
    ...  
}  
  
void foo_listener() {  
    while(1) {  
        recv, call foo  
    }  
}
```

# RPC

## Machine A

```
int main(...) {  
    int x = foo("hello");  
}
```

client  
wrapper

```
int foo(char *msg) {  
    send msg to B  
    recv msg from B  
}
```

## Machine B

```
int foo(char *msg) {  
    ...  
}
```

server  
wrapper

```
void foo_listener() {  
    while(1) {  
        recv, call foo  
    }  
}
```

# RPC TOOLS

RPC packages help with two components

## (1) Runtime library

- Thread pool
- Socket listeners call functions on server

## (2) Stub generation

- Create wrappers automatically
- Many tools available (rpcgen, thrift, protobufs)

# WRAPPER GENERATION

Wrappers must do conversions:

- client arguments to message
- message to server arguments
- convert server return value to message
- convert message to client return value

Need uniform endianness (wrappers do this)

Conversion is called marshaling/unmarshaling, or serializing/deserializing




# WRAPPER GENERATION: POINTERS

Why are pointers problematic?

Address passed from client not valid on server

Solutions? Smart RPC package: follow pointers and copy data

# RPC OVER TCP?




# RPC OVER UDP

Strategy: use function return as implicit ACK

Piggybacking technique

What if function takes a long time?  
then send a separate ACK



## Other Issues

- **Long-running calls**, client periodically asks server for results
- **Data Organization** – e.g. Big-Endian vs. Little Endian
  - Sun's XDR (eXternal Data Representation) formatting standard
  - Google's gRDP uses HTTP/2
- Some systems provide both **synchronous** (i.e. wait for result) and **asynchronous** (i.e. return immediately with some type of callback)

# Summary

- UDP for unreliable communication
- TCP for reliable communication
- RPC often builds on top of UDP layer, handles communication failures itself
  - has a stub generator and run-time library
  - handles issues like fragmentation and byte ordering
  - Typically synchronous calls (wait for completion)
- RPC packages include:
  - Sun's RPC system
  - Google's gRPC
  - Apache Thrift
  - JSON-RPC